

```
import torch
from torch import nn
from torch import optim
import numpy as np
import data
import matplotlib.pyplot as plt
from typing import List, Callable
from sklearn.model_selection import train_test_split
from copy import deepcopy
import math
import time
```

```
class PTDeep(nn.Module):
```

```
    def __init__(self, dimensions: List[int], activations: List[Callable], param_lambda=0.0, use_cuda=False):
        super(PTDeep, self).__init__()
        device = "cuda" if use_cuda is True else "cpu"
        D = dimensions[0]
        C = dimensions[-1]
        weights = []
        biases = []
        for i in range(len(dimensions) - 1):
            dim_in = dimensions[i]
            dim_out = dimensions[i+1]
            W = np.random.randn(dim_in, dim_out).astype(np.float64)
            b = np.random.randn(1, dim_out).astype(np.float64)
            weights.append(nn.Parameter(data=torch.from_numpy(W).to(device=device), requires_grad=True))
            biases.append(nn.Parameter(data=torch.from_numpy(b).to(device=device), requires_grad=True))
```

```
        self.weights = nn.ParameterList(weights)
        self.biases = nn.ParameterList(biases)
        self.activations = activations
        self.param_lambda = param_lambda
        self.loss_trace = None
```

```
    def forward(self, X):
        h = X
        for i in range(len(self.weights)):
            scores = torch.mm(h, self.weights[i]) + self.biases[i]
            h = self.activations[i](scores)

        return h
```

```
    def get_loss(self, X, Yoh_):
        N = len(X)
        probs = self.forward(X)
        logprobs = torch.log(probs)
        loss = - 1 / N * torch.sum(logprobs * Yoh_)
        return loss
```

```
    def train(self, X, Yoh_, param_niter=10_000, param_delta=0.05, print_frequency=200,
               epsilon=1e-3, trace=True, early_stopping=False):
        optimizer = optim.SGD(self.parameters(), lr=param_delta, weight_decay=self.param_lambda)
        best_validation_loss = None
        if trace is True:
            self.loss_trace = []
        if early_stopping is True:
            best_weights = None
            best_biases = None
            x_train, x_validate, y_train, y_validate = train_test_split(X, Yoh_, test_size=0.2)
```

```
        start = time.time()
        for i in range(param_niter):
            # calculate loss
            loss = self.get_loss(x_train, y_train)
            # if trace is true, remember this loss
```

```

if trace is True:
    self.loss_trace.append(loss.detach().numpy())
# periodically print status
if i % print_frequency == 0:
    end = time.time()
    print(f"iteration {i} loss {loss} duration {end-start}")
# calculate the gradients and adjust the parameters
loss.backward()
optimizer.step()
optimizer.zero_grad()
# if early stopping is on, validate and remember best models parameters
if early_stopping is True:
    validation_loss = self.get_loss(x_validate, y_validate)
    optimizer.zero_grad()
    if best_validation_loss is None or validation_loss < best_validation_loss:
        best_validation_loss = validation_loss
        best_weights = deepcopy(self.weights)
        best_biases = deepcopy(self.biases)

if early_stopping is True:
    self.weights = best_weights
    self.biases = best_biases

return self

def train_mb(self, X, Yoh_,
             param_niter=10_000, param_delta=0.05, print_frequency=200, batch_size=128,
             trace=True, early_stopping=False, optimizer=None, scheduler=None):
    if optimizer is None:
        optimizer = optim.SGD(self.parameters(), lr=param_delta, weight_decay=self.param_lambda)

    best_validation_loss = None
    if trace is True:
        self.loss_trace = []
    if early_stopping is True:
        best_weights = None
        best_biases = None
        x_train, x_validate, y_train, y_validate = train_test_split(X, Yoh_, test_size=0.2)
    else:
        x_train, y_train = X, Yoh_

    start = time.time()
    for i in range(param_niter):
        # shuffle and create batches
        p = np.random.permutation(len(x_train))
        x_train = x_train[p]
        y_train = y_train[p]
        number_of_batches = int(math.ceil(len(x_train) / batch_size))
        # iterate over all batches
        for j in range(number_of_batches):
            x_batch = x_train[(j*batch_size):(j+1)*batch_size]
            y_batch = y_train[(j*batch_size):(j+1)*batch_size]
            loss = self.get_loss(x_batch, y_batch)
            # if trace is true, remember this loss
            if trace is True:
                self.loss_trace.append(loss.detach().numpy())
            # periodically print status
            if i % print_frequency == 0 and j % print_frequency == 0:
                end = time.time()
                print(f"iteration {i} batch {j} loss {loss} duration {end-start}")
            # calculate the gradients and adjust the parameters
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()

        # if scheduler is given, step it up

```

```

if scheduler is not None:
    scheduler.step()
# if early stopping is on, validate and remember best models parameters
if early_stopping is True:
    validation_loss = self.get_loss(x_validate, y_validate)
    optimizer.zero_grad()
    if best_validation_loss is None or validation_loss < best_validation_loss:
        best_validation_loss = validation_loss
    elif validation_loss > best_validation_loss:
        print("Validation loss is going up! Ending training.")
        break

return self

def eval(self, X):
    X = torch.from_numpy(X)
    probs = self.forward(X)
    return probs.detach().numpy()

def my_softmax(X):
    return torch.softmax(X, 1)

def count_params(model: nn.Module):
    count = 0
    for name, param in model.named_parameters():
        print(f"{name}: {param.shape}")
        count += param.shape[0] * param.shape[1]
    print("number of params:", count)

if __name__ == "__main__":
    np.random.seed(101)
    X, Y_ = data.sample_gmm_2d(4, 2, 40)
    Yoh_ = data.class_to_onehot(Y_)

    model = PTDeep([2, 10, 2], [torch.sigmoid, my_softmax])
    model.train(torch.from_numpy(X), torch.from_numpy(Yoh_), 10_000, 0.05, 200)
    probs = model.eval(X)
    # predicted classes
    Y = np.hstack([np.argmax(probs[i][:],) for i in range(probs.shape[0])])
    # reshaping for other methods purposes
    Y_ = np.hstack(Y_)
    accuracy, pr, M = data.eval_perf_multi(Y, Y_)
    print("Accuracy: ", accuracy)
    print("Precision / Recall: ", pr)
    print("Confussion Matrix: ", M)
    print("All parameters:")
    count_params(model)

def decfun(X):
    X = torch.from_numpy(X)
    to_return = model.forward(X)
    return to_return.detach().numpy().argmax(axis=1)

# decfun za binlogreg
decfun = lambda X: model.forward(torch.from_numpy(X)).detach().numpy()[:, 1]
bbox = (np.min(X, axis=0), np.max(X, axis=0))
data.graph_surface(decfun, bbox, offset=0.5)
# graph the data points
data.graph_data(X, Y_, Y, special=[])
# show the plot
plt.show()

```